

Programming Homework 3 - Report

1. Parallel Matrix Multiplication

1.1 Naive MatMul with OpenMP

| No. of Threads | Execution Time | Performance |
|----------------|-----------------|-------------------|
| 1 | 1366.138831 sec | 100.603943 MFLOPS |
| 4 | 375.451539 sec | 366.063098 MFLOPS |
| 16 | 394.123210 sec | 348.720781 MFLOPS |
| 64 | 390.727256 sec | 351.751641 MFLOPS |
| 256 | 386.786432 sec | 355.335508 MFLOPS |

1.2 Block MatMul with OpenMP

| No. of Threads | Execution Time | Performance |
|----------------|----------------|--------------------|
| 1 | 319.371676 sec | 430.341711 MFLOPS |
| 4 | 101.459692 sec | 1354.616314 MFLOPS |
| 16 | 95.469828 sec | 1439.606168 MFLOPS |
| 64 | 101.700691 sec | 1351.406295 MFLOPS |
| 256 | 90.865101 sec | 1512.560400 MFLOPS |


1.3 Transposed MatMul with OpenMP

| No. of Threads | Execution Time | Performance |
|----------------|----------------|--------------------|
| 1 | 277.867787 sec | 494.619960 MFLOPS |
| 4 | 86.697244 sec | 1585.274763 MFLOPS |
| 16 | 91.441576 sec | 1503.024768 MFLOPS |
| 64 | 91.454943 sec | 502.805084 MFLOPS |
| 256 | 77.599605 sec | 1771.129540 MFLOPS |

Sample Output

```
Number of FLOPs = 137438953472, Execution time = 390.727256 sec,  
351.751641 MFLOPs per sec  
C[100][100]=879616000.000000
```

Execution Time vs No. of Threads



Observations:

1. The execution time of naive matrix multiplication reduces with the number of threads as expected. However, it plateaus quickly because it is being run on a 4-core machine, and there's not much room left to improve performance with more threads.

2. The execution times for Block and Transposed Matrix Multiplication are much lower compared to the Naive Matrix Multiplication because both block and transposed Matrix Multiplication are cache-optimized methods, wherein the next element in the loop is usually found in the cache. In contrast, the naive matrix multiplication iterates column-wise on a row-major matrix.

2. Estimating Pi

Serial time: 0.243284 sec

Version a: DO/for OpenMP parallel with 4 threads: 0.100009 sec

Version b: 2 Threads with SECTIONS: 0.122512 sec

Observation: We observe a significant speed-up in both versions compared to the serial time, especially when using the REDUCTION data attribute.

Output:

```
> ./p2b  
Estimated pi is 3.142402, execution time = 0.134847 sec
```

3. Quick Sort

3.1 Serial Implementation: 7663.608671 sec

Note: I had to increase my system stack size using “ulimit -s unlimited” to fix a segmentation fault caused by stack overflow.

3.2 Sequential initialization, OpenMP parallel, 2 threads: 7225.312135 sec

Note: Here, OpenMP’s child processes don’t have the same stack size as the master thread, and to increase the stack size of the child processes, I set an env variable called OMP_STACKSIZE. Without this, I was facing stack overflow.

3.3 Random initialization, OpenMP parallel, 2 threads: 0.294041 sec

Observation:

1. We do not observe a significant time difference between the serial and OpenMP parallel implementation. This is because, in both cases, the stack memory (which is in RAM) acts as the bottleneck, and accessing RAM using one thread or two threads won't make a meaningful difference as the unit delay will be similar.
2. The random initialization is significantly faster because the quicksort algorithm has a worst-case time complexity of $O(n^2)$ when the array is sorted in reverse order (Case 3.1 and 3.2). In contrast, the average case time complexity is $O(n \log n)$ which is when the array is in no particular order (Case 3.3). The significant performance gain in 3.3 can also be attributed to the recursion stack not going as deep, avoiding the RAM access delay by many folds.

Output:

```
> ./p3c
122 2066 2073 3406 5536 6518 6565 8520 9121 11895 12170 13050 14097 14195 14385 14939
Execution time = 0.294041 sec
```